

ASSIGNMENT11.2

Perform incremental load in Hive
Read from MySQL Table and load it in Hive table.
Create hive table if it does not exist.
If it exists, perform the incremental load.

```
mysql> create table employee1
-> (id int,name varchar(20),salary int);
Query OK, 0 rows affected (0.07 sec)
```

```
mysql> insert into employee1 values(1,'vishnu',500);
Query OK, 1 row affected (0.04 sec)
```

```
mysql> insert into employee1 values(2,'vipin',300);
Query OK, 1 row affected (0.00 sec)
```

```
[acadgild@localhost ~]$ sqoop import --connect
jdbc:mysql://localhost/db1 --username 'root' -P --table 'employee1'
--target-dir '/display' --hive-import --create-hive-table
--hive-table 'employeehive' -m 1
Warning: /usr/local/sqoop/../../hcatalog does not exist! HCatalog jobs
will fail.
Please set $HCAT_HOME to the root of your HCatalog installation.
Warning: /usr/local/sqoop/../../accumulo does not exist! Accumulo
imports will fail.
Please set $ACCUMULO_HOME to the root of your Accumulo installation.
Warning: /usr/local/sqoop/../../zookeeper does not exist! Accumulo
imports will fail.
Please set $ZOOKEEPER_HOME to the root of your Zookeeper
installation.
2017-11-03 23:32:07,098 INFO [main] sqoop.Sqoop: Running Sqoop
version: 1.4.5
Enter password:
2017-11-03 23:32:08,944 INFO [main] tool.BaseSqoopTool: Using
Hive-specific delimiters for output. You can override
2017-11-03 23:32:08,944 INFO [main] tool.BaseSqoopTool: delimiters
with --fields-terminated-by, etc.
2017-11-03 23:32:09,463 INFO [main] manager.MySQLManager: Preparing
to use a MySQL streaming resultset.
2017-11-03 23:32:09,491 INFO [main] tool.CodeGenTool: Beginning code
generation
2017-11-03 23:32:10,162 INFO [main] manager.SqlManager: Executing
SQL statement: SELECT t.* FROM `employee` AS t LIMIT 1
2017-11-03 23:32:10,271 INFO [main] manager.SqlManager: Executing
SQL statement: SELECT t.* FROM `employee` AS t LIMIT 1
2017-11-03 23:32:10,294 INFO [main] orm.CompilationManager:
HADOOP_MAPRED_HOME is /usr/local/hadoop-2.6.0
```

Note:

/tmp/sqoop-acadgild/compile/77e0717fd2a983e81d99806d7310dfea/employee.java uses or overrides a deprecated API.

Note: Recompile with -Xlint:deprecation for details.

2017-11-03 23:32:14,419 INFO [main] orm.CompilationManager: Writing jar file:

/tmp/sqoop-acadgild/compile/77e0717fd2a983e81d99806d7310dfea/employee.jar

2017-11-03 23:32:14,468 WARN [main] manager.MySQLManager: It looks like you are importing from mysql.

2017-11-03 23:32:14,468 WARN [main] manager.MySQLManager: This transfer can be faster! Use the --direct

2017-11-03 23:32:14,468 WARN [main] manager.MySQLManager: option to exercise a MySQL-specific fast path.

2017-11-03 23:32:14,468 INFO [main] manager.MySQLManager: Setting zero DATETIME behavior to convertToNull (mysql)

2017-11-03 23:32:14,483 INFO [main] mapreduce.ImportJobBase: Beginning import of employee

SLF4J: Class path contains multiple SLF4J bindings.

SLF4J: Found binding in

[jar:file:/usr/local/hbase/lib/slf4j-log4j12-1.6.4.jar!/org/slf4j/impl/StaticLoggerBinder.class]

SLF4J: Found binding in

[jar:file:/usr/local/hadoop-2.6.0/share/hadoop/common/lib/slf4j-log4j12-1.7.5.jar!/org/slf4j/impl/StaticLoggerBinder.class]

SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.

2017-11-03 23:32:15,175 WARN [main] util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

2017-11-03 23:32:15,198 INFO [main] Configuration.deprecation: mapred.jar is deprecated. Instead, use mapreduce.job.jar

2017-11-03 23:32:16,450 INFO [main] Configuration.deprecation: mapred.map.tasks is deprecated. Instead, use mapreduce.job.maps

2017-11-03 23:32:16,844 INFO [main] client.RMProxy: Connecting to ResourceManager at /0.0.0.0:8032

2017-11-03 23:32:21,565 INFO [main] db.DBInputFormat: Using read committed transaction isolation

2017-11-03 23:32:21,766 INFO [main] mapreduce.JobSubmitter: number of splits:1

2017-11-03 23:32:21,844 INFO [main] Configuration.deprecation: mapred.job.name is deprecated. Instead, use mapreduce.job.name

2017-11-03 23:32:21,845 INFO [main] Configuration.deprecation: mapred.cache.files.timestamps is deprecated. Instead, use mapreduce.job.cache.files.timestamps

2017-11-03 23:32:21,845 INFO [main] Configuration.deprecation: mapreduce.map.class is deprecated. Instead, use mapreduce.job.map.class

2017-11-03 23:32:21,846 INFO [main] Configuration.deprecation:

```
mapreduce.inputformat.class is deprecated. Instead, use
mapreduce.job.inputformat.class
2017-11-03 23:32:21,848 INFO [main] Configuration.deprecation:
mapreduce.outputformat.class is deprecated. Instead, use
mapreduce.job.outputformat.class
2017-11-03 23:32:21,848 INFO [main] Configuration.deprecation:
mapred.output.value.class is deprecated. Instead, use
mapreduce.job.output.value.class
2017-11-03 23:32:21,849 INFO [main] Configuration.deprecation:
mapred.output.dir is deprecated. Instead, use
mapreduce.output.fileoutputformat.outputdir
2017-11-03 23:32:21,849 INFO [main] Configuration.deprecation:
mapred.cache.files is deprecated. Instead, use
mapreduce.job.cache.files
2017-11-03 23:32:21,849 INFO [main] Configuration.deprecation:
mapred.working.dir is deprecated. Instead, use
mapreduce.job.working.dir
2017-11-03 23:32:21,849 INFO [main] Configuration.deprecation:
mapred.job.classpath.files is deprecated. Instead, use
mapreduce.job.classpath.files
2017-11-03 23:32:21,849 INFO [main] Configuration.deprecation:
user.name is deprecated. Instead, use mapreduce.job.user.name
2017-11-03 23:32:21,849 INFO [main] Configuration.deprecation:
mapred.reduce.tasks is deprecated. Instead, use mapreduce.job.reduces
2017-11-03 23:32:21,849 INFO [main] Configuration.deprecation:
mapred.cache.files.filesizes is deprecated. Instead, use
mapreduce.job.cache.files.filesizes
2017-11-03 23:32:21,849 INFO [main] Configuration.deprecation:
mapred.output.key.class is deprecated. Instead, use
mapreduce.job.output.key.class
2017-11-03 23:32:22,368 INFO [main] mapreduce.JobSubmitter:
Submitting tokens for job: job_1509729939048_0001
2017-11-03 23:32:23,521 INFO [main] impl.YarnClientImpl: Submitted
application application_1509729939048_0001 to ResourceManager at /
0.0.0.0:8032
2017-11-03 23:32:23,690 INFO [main] mapreduce.Job: The url to track
the job:
http://http://localhost:8088/proxy/application_1509729939048_0001/
2017-11-03 23:32:23,691 INFO [main] mapreduce.Job: Running job:
job_1509729939048_0001
2017-11-03 23:32:46,419 INFO [main] mapreduce.Job: Job
job_1509729939048_0001 running in uber mode : false
2017-11-03 23:32:46,429 INFO [main] mapreduce.Job: map 0% reduce 0%
2017-11-03 23:32:58,765 INFO [main] mapreduce.Job: map 100% reduce
0%
2017-11-03 23:32:59,834 INFO [main] mapreduce.Job: Job
job_1509729939048_0001 completed successfully
2017-11-03 23:33:00,215 ERROR [main] tool.ImportTool: Imported
Failed: No enum constant
```

org.apache.hadoop.mapreduce.JobCounter.MB_MILLIS_MAPS

```
[acadgild@localhost ~]$ hadoop fs -cat /display/part-m-00000
17/11/03 23:35:21 WARN util.NativeCodeLoader: Unable to load
native-hadoop library for your platform... using builtin-java classes
where applicable
1#vishnu#500
2#vipin#300
```

Incremental importing

```
mysql> insert into employee1 values(3,'vishnuvijay',500);
Query OK, 1 row affected (0.00 sec)
```

```
mysql> insert into employee1 values(4,'vijay',500);
Query OK, 1 row affected (0.00 sec)
```

```
[acadgild@localhost ~]$ sqoop import --connect
jdbc:mysql://localhost/db1 --username 'root' -P --table 'employee1'
--target-dir '/display' --hive-import --create-hive-table
--hive-table 'employeehive1' --incremental append --check-column id
--last-value 3 -m 1
Warning: /usr/local/sqoop/./hcatalog does not exist! HCatalog jobs
will fail.
Please set $HCAT_HOME to the root of your HCatalog installation.
Warning: /usr/local/sqoop/./accumulo does not exist! Accumulo
imports will fail.
Please set $ACCUMULO_HOME to the root of your Accumulo installation.
Warning: /usr/local/sqoop/./zookeeper does not exist! Accumulo
imports will fail.
Please set $ZOOKEEPER_HOME to the root of your Zookeeper
installation.
2017-11-04 01:09:52,322 INFO [main] sqoop.Sqoop: Running Sqoop
version: 1.4.5
Enter password:
2017-11-04 01:09:54,197 INFO [main] tool.BaseSqoopTool: Using
Hive-specific delimiters for output. You can override
2017-11-04 01:09:54,197 INFO [main] tool.BaseSqoopTool: delimiters
with --fields-terminated-by, etc.
2017-11-04 01:09:54,791 INFO [main] manager.MySQLManager: Preparing
to use a MySQL streaming resultset.
2017-11-04 01:09:54,813 INFO [main] tool.CodeGenTool: Beginning code
generation
2017-11-04 01:09:55,534 INFO [main] manager.SqlManager: Executing
SQL statement: SELECT t.* FROM `employee1` AS t LIMIT 1
2017-11-04 01:09:55,669 INFO [main] manager.SqlManager: Executing
SQL statement: SELECT t.* FROM `employee1` AS t LIMIT 1
2017-11-04 01:09:55,714 INFO [main] orm.CompilationManager:
```

HADOOP_MAPRED_HOME is /usr/local/hadoop-2.6.0

Note:

/tmp/sqoop-acadgild/compile/ce8dd9bd0a6f4247b7a683c483db3490/employee1.java uses or overrides a deprecated API.

Note: Recompile with -Xlint:deprecation for details.

2017-11-04 01:10:01,941 INFO [main] orm.CompilationManager: Writing jar file:

/tmp/sqoop-acadgild/compile/ce8dd9bd0a6f4247b7a683c483db3490/employee1.jar

SLF4J: Class path contains multiple SLF4J bindings.

SLF4J: Found binding in

[jar:file:/usr/local/hbase/lib/slf4j-log4j12-1.6.4.jar!/org/slf4j/impl/StaticLoggerBinder.class]

SLF4J: Found binding in

[jar:file:/usr/local/hadoop-2.6.0/share/hadoop/common/lib/slf4j-log4j12-1.7.5.jar!/org/slf4j/impl/StaticLoggerBinder.class]

SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.

2017-11-04 01:10:03,211 WARN [main] util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

2017-11-04 01:10:04,687 INFO [main] tool.ImportTool: Maximal id query for free form incremental import: SELECT MAX(`id`) FROM employee1

2017-11-04 01:10:04,689 INFO [main] tool.ImportTool: Incremental import based on column `id`

2017-11-04 01:10:04,689 INFO [main] tool.ImportTool: Lower bound value: 3

2017-11-04 01:10:04,689 INFO [main] tool.ImportTool: Upper bound value: 4

2017-11-04 01:10:04,690 WARN [main] manager.MySQLManager: It looks like you are importing from mysql.

2017-11-04 01:10:04,690 WARN [main] manager.MySQLManager: This transfer can be faster! Use the --direct

2017-11-04 01:10:04,690 WARN [main] manager.MySQLManager: option to exercise a MySQL-specific fast path.

2017-11-04 01:10:04,690 INFO [main] manager.MySQLManager: Setting zero DATETIME behavior to convertToNull (mysql)

2017-11-04 01:10:04,698 INFO [main] mapreduce.ImportJobBase: Beginning import of employee1

2017-11-04 01:10:04,753 INFO [main] Configuration.deprecation: mapred.jar is deprecated. Instead, use mapreduce.job.jar

2017-11-04 01:10:04,847 INFO [main] Configuration.deprecation: mapred.map.tasks is deprecated. Instead, use mapreduce.job.maps

2017-11-04 01:10:05,393 INFO [main] client.RMProxy: Connecting to ResourceManager at /0.0.0.0:8032

2017-11-04 01:10:08,491 INFO [main] db.DBInputFormat: Using read committed transaction isolation

2017-11-04 01:10:08,622 INFO [main] mapreduce.JobSubmitter: number

```
of splits:1
2017-11-04 01:10:08,662 INFO [main] Configuration.deprecation:
mapred.job.name is deprecated. Instead, use mapreduce.job.name
2017-11-04 01:10:08,664 INFO [main] Configuration.deprecation:
mapred.cache.files.timestamps is deprecated. Instead, use
mapreduce.job.cache.files.timestamps
2017-11-04 01:10:08,664 INFO [main] Configuration.deprecation:
mapreduce.map.class is deprecated. Instead, use
mapreduce.job.map.class
2017-11-04 01:10:08,667 INFO [main] Configuration.deprecation:
mapreduce.inputformat.class is deprecated. Instead, use
mapreduce.job.inputformat.class
2017-11-04 01:10:08,672 INFO [main] Configuration.deprecation:
mapreduce.outputformat.class is deprecated. Instead, use
mapreduce.job.outputformat.class
2017-11-04 01:10:08,672 INFO [main] Configuration.deprecation:
mapred.output.value.class is deprecated. Instead, use
mapreduce.job.output.value.class
2017-11-04 01:10:08,673 INFO [main] Configuration.deprecation:
mapred.output.dir is deprecated. Instead, use
mapreduce.output.fileoutputformat.outputdir
2017-11-04 01:10:08,674 INFO [main] Configuration.deprecation:
mapred.cache.files is deprecated. Instead, use
mapreduce.job.cache.files
2017-11-04 01:10:08,674 INFO [main] Configuration.deprecation:
mapred.working.dir is deprecated. Instead, use
mapreduce.job.working.dir
2017-11-04 01:10:08,674 INFO [main] Configuration.deprecation:
mapred.job.classpath.files is deprecated. Instead, use
mapreduce.job.classpath.files
2017-11-04 01:10:08,675 INFO [main] Configuration.deprecation:
user.name is deprecated. Instead, use mapreduce.job.user.name
2017-11-04 01:10:08,675 INFO [main] Configuration.deprecation:
mapred.reduce.tasks is deprecated. Instead, use mapreduce.job.reduces
2017-11-04 01:10:08,675 INFO [main] Configuration.deprecation:
mapred.cache.files.filesizes is deprecated. Instead, use
mapreduce.job.cache.files.filesizes
2017-11-04 01:10:08,675 INFO [main] Configuration.deprecation:
mapred.output.key.class is deprecated. Instead, use
mapreduce.job.output.key.class
2017-11-04 01:10:09,038 INFO [main] mapreduce.JobSubmitter:
Submitting tokens for job: job_1509729939048_0014
2017-11-04 01:10:09,978 INFO [main] impl.YarnClientImpl: Submitted
application application_1509729939048_0014 to ResourceManager at /
0.0.0.0:8032
2017-11-04 01:10:10,211 INFO [main] mapreduce.Job: The url to track
the job:
http://http://localhost:8088/proxy/application_1509729939048_0014/
2017-11-04 01:10:10,233 INFO [main] mapreduce.Job: Running job:
```

```
job_1509729939048_0014
2017-11-04 01:10:27,005 INFO [main] mapreduce.Job: Job
job_1509729939048_0014 running in uber mode : false
2017-11-04 01:10:27,026 INFO [main] mapreduce.Job: map 0% reduce 0%
2017-11-04 01:10:43,703 INFO [main] mapreduce.Job: map 100% reduce
0%
2017-11-04 01:10:43,750 INFO [main] mapreduce.Job: Job
job_1509729939048_0014 completed successfully
2017-11-04 01:10:44,204 ERROR [main] tool.ImportTool: Imported
Failed: No enum constant
org.apache.hadoop.mapreduce.JobCounter.MB_MILLIS_MAPS
```

```
[acadgild@localhost ~]$ hadoop fs -cat /display/part-m-0000017/11/04
01:17:46 WARN util.NativeCodeLoader: Unable to load native-hadoop
library for your platform... using builtin-java classes where
applicable
3#vishnuvijay#500
4#vijay#500
```